

Sequence Labeling

CS4742 Natural Language Processing

Lecture 08

Jiho Noh

Department of Computer Science
Kennesaw State University

CS4742 Summer 2025 ¹

¹This lecture is based on the slides from Dr. Hafiz Khan at KSU.

1 Problem Formalization

2 Part-of-Speech (POS) Tagging

3 Hidden Markov Model (HMM)

4 Named Entity Recognition (NER)

Sequence Labeling

- Sequence labeling assigns a label $y \in Y$ to each element in a given sequence x .
- If the input is a text sequence, elements are words or characters:
 $x = \{w_1, w_2, \dots, w_n\}$, where n denotes the length of the text input.
- The objective is to find the best label sequence $y = \{y_1, y_2, \dots, y_n\}$ for the input sequence x .
- **Algorithms:** Hidden Markov Model (HMM), Conditional Random Field (CRF), Recurrent Neural Network (RNN).

Sequence Labeling Tasks

- Y can be any set of labels. The most common tasks are:
 - ▶ **Parts of Speech (POS)** - noun, verb, pronoun, preposition, adverb, conjunction, participle, and article.
 - ▶ **Name Entity Recognition** - person, organization, location, date, time, money, etc.
- **Parts of Speech** and **Name entities** are useful clues for understanding sentence structure and meaning.
- Both are crucial for extracting meaningful information and improving the accuracy of NLP applications.

- 1 Problem Formalization
- 2 Part-of-Speech (POS) Tagging**
- 3 Hidden Markov Model (HMM)
- 4 Named Entity Recognition (NER)

Part-of-Speech Tagging (POS Tagging)

POS tagging is the process of assigning a part of speech to each word in a sentence. The part of speech can be *noun*, *verb*, *adjective*, *adverb*, etc. A set of all POS tags is called the **tagset**, and various tagsets exist.

POS Tagging example:

- (8.1) There/**PRO**/**EX** are/**VERB**/**VBP** 70/**NUM**/**CD** children/**NOUN**/**NNS**
there/**ADV**/**RB** ./**PUNC**/.
- (8.2) Preliminary/**ADJ**/**JJ** findings/**NOUN**/**NNS** were/**AUX**/**VBD** reported/**VERB**/**VBN**
in/**ADP**/**IN** today/**NOUN**/**NN** 's/**PART**/**POS** New/**PROPN**/**NNP**
England/**PROPN**/**NNP** Journal/**PROPN**/**NNP** of/**ADP**/**IN** Medicine/**PROPN**/**NNP**

Universal Dependencies (UD) Tag Set

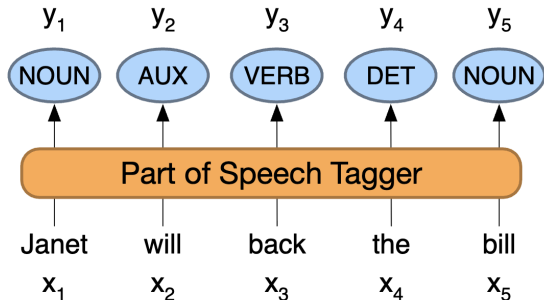
	Tag	Description	Example
Open Class	ADJ	Adjective: noun modifiers describing properties	<i>red, young, awesome</i>
	ADV	Adverb: verb modifiers of time, place, manner	<i>very, slowly, home, yesterday</i>
	NOUN	words for persons, places, things, etc.	<i>algorithm, cat, mango, beauty</i>
	VERB	words for actions and processes	<i>draw, provide, go</i>
	PROPN	Proper noun: name of a person, organization, place, etc..	<i>Regina, IBM, Colorado</i>
	INTJ	Interjection: exclamation, greeting, yes/no response, etc.	<i>oh, um, yes, hello</i>
Closed Class Words	ADP	Adposition (Preposition/Postposition): marks a noun's spacial, temporal, or other relation	<i>in, on, by, under</i>
	AUX	Auxiliary: helping verb marking tense, aspect, mood, etc.,	<i>can, may, should, are</i>
	CCONJ	Coordinating Conjunction: joins two phrases/clauses	<i>and, or, but</i>
	DET	Determiner: marks noun phrase properties	<i>a, an, the, this</i>
	NUM	Numeral	<i>one, two, first, second</i>
	PART	Particle: a preposition-like form used together with a verb	<i>up, down, on, off, in, out, at, by</i>
	PRON	Pronoun: a shorthand for referring to an entity or event	<i>she, who, I, others</i>
Other	SCONJ	Subordinating Conjunction: joins a main clause with a subordinate clause such as a sentential complement	<i>that, which</i>
	PUNCT	Punctuation	<i>; , ()</i>
	SYM	Symbols like \$ or emoji	<i>\$, %</i>
	X	Other	<i>asdf, qwfg</i>

Penn Treebank Tag Set

Tag	Description	Example	Tag	Description	Example	Tag	Description	Example
CC	coord. conj.	<i>and, but, or</i>	NNP	proper noun, sing.	<i>IBM</i>	TO	“to”	<i>to</i>
CD	cardinal number	<i>one, two</i>	NNPS	proper noun, plu.	<i>Carolinas</i>	UH	interjection	<i>ah, oops</i>
DT	determiner	<i>a, the</i>	NNS	noun, plural	<i>llamas</i>	VB	verb base	<i>eat</i>
EX	existential ‘there’	<i>there</i>	PDT	predeterminer	<i>all, both</i>	VBD	verb past tense	<i>ate</i>
FW	foreign word	<i>mea culpa</i>	POS	possessive ending	<i>’s</i>	VBG	verb gerund	<i>eating</i>
IN	preposition/ subordin-conj	<i>of, in, by</i>	PRP	personal pronoun	<i>I, you, he</i>	VBN	verb past partici- ple	<i>eaten</i>
JJ	adjective	<i>yellow</i>	PRP\$	possess. pronoun	<i>your, one’s</i>	VBP	verb non-3sg-pr	<i>eat</i>
JJR	comparative adj	<i>bigger</i>	RB	adverb	<i>quickly</i>	VBZ	verb 3sg pres	<i>eats</i>
JJS	superlative adj	<i>wildest</i>	RBR	comparative adv	<i>faster</i>	WDT	wh-determ.	<i>which, that</i>
LS	list item marker	<i>1, 2, One</i>	RBS	superlatv. adv	<i>fastest</i>	WP	wh-pronoun	<i>what, who</i>
MD	modal	<i>can, should</i>	RP	particle	<i>up, off</i>	WP\$	wh-possess.	<i>whose</i>
NN	sing or mass noun	<i>llama</i>	SYM	symbol	<i>+, %, &</i>	WRB	wh-adverb	<i>how, where</i>

POS Tagging

- **Tagging** is a *disambiguation* task; words are ambiguous.
- A word can have more than one *ambiguous* possible part-of-speech.
- The goal is to find the **correct** tag for the context.



Ambiguity Resolution

6 different parts of speech for the word *back*

- earnings growth took a back/**JJ** seat
- a small building in the back/**NN**
- a clear majority of senators back/**VBP** the bill
- Dave began to back/**VB** toward the door
- enable the country to buy back/**RP** debt
- I was twenty-one back/**RB** then

Parts of Speech Tagging Performance

- The *accuracy* of part-of-speech tagging algorithms (the percentage of test set tags that match human gold labels) is extremely high.
- POS Tagging is known to be a **solved problem** in NLP.
- However, POS tagging can be found in various tasks:
 - ▶ Information retrieval, parsing, Text to Speech (TTS) applications, information extraction, linguistic research for corpora
 - ▶ POS can be used as an intermediate step for higher level NLP tasks – parsing, semantics analysis, translation etc.

Methods — First Take

Baseline: Maximum Likelihood

- Given an *ambiguous word*, choose the tag which is most frequent in the training corpus.
- $\arg \max_{\text{tag}} P(\text{tag}|\text{w}) = \arg \max_{\text{tag}} \frac{P(\text{tag}, \text{w})}{P(\text{w})} = \arg \max_{\text{tag}} P(\text{tag}, \text{w})$

Better Methods

Rule-based POS tagging:

- A set of handwritten rules and use *contextual information* to assign POS tags to words.
- **Example Rule:** if an ambiguous/unknown word ends with the suffix “ing” and is preceded by a word likely to be verb, tag it as a verb.
 - ▶ He is playing football.
 - ▶ “Play” – tag as “**Verb**”

Transformation based tagging:

- A pre-defined handcrafted rules combined with automatically rules.
- Automatic rules are generated from training data.

Even Better Methods

- **Deep learning models:** various deep models to tag *POS*
- **Probabilistic (stochastic) tagging:**
 - ▶ A stochastic approach computes frequency, probability from training set. During testing select highest probable tag.
 - ▶ Example algorithm: **Hidden Markov Model**

- 1 Problem Formalization
- 2 Part-of-Speech (POS) Tagging
- 3 Hidden Markov Model (HMM)**
- 4 Named Entity Recognition (NER)

2

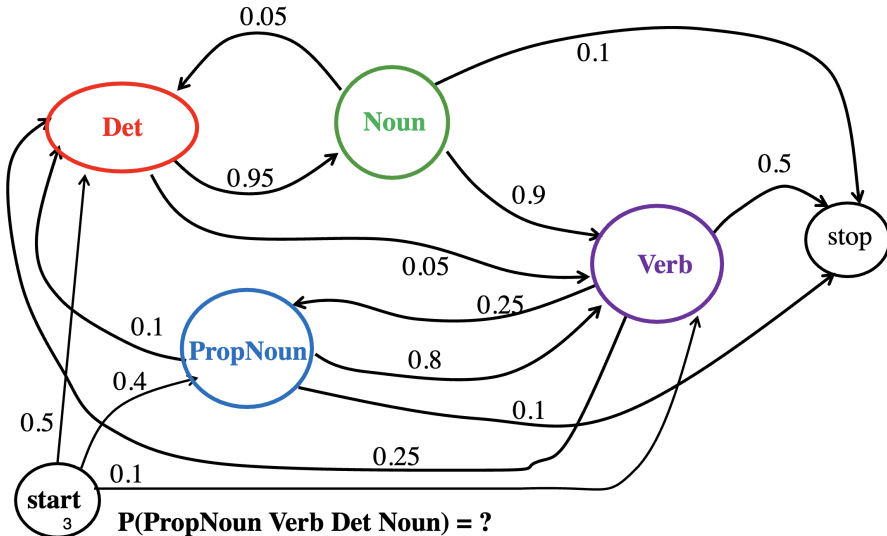
²Following slides are modified from Prof. Claire Cardie's slides and Prof. Raymond Mooney's slides. Some of the graphs are taken from the textbook.

POS Tagging using a Markov Model (Markov Chain)


- **Random variables** are used to represent the sequence of part-of-speech tags assigned to words in a sentence.
- With this model, each random variable corresponds to a word and takes on a value from the set of possible *POS tags*.
- A sequence of random variables X_1, X_2, \dots, X_n is said to be a **Markov chain** if the conditional probability of the next state depends only on the *current state* and not on the previous states (**Markov Assumption**).

$$P(X_{n+1} = x | X_n = x_n, X_{n-1} = x_{n-1}, \dots, X_1 = x_1) = P(X_{n+1} = x | X_n = x_n)$$

Sample Markov Model for POS



Sample Markov Model for POS



Hidden Markov Model for POS Tagging

What do we know?


- **States** = POS tags
- **Transitions Probabilities** = probabilities of moving from one state to another

What do we not see in the previous figure?

- **Observation** = a sequence of words
- **Emission Probabilities** = probabilities of observing a word given a state (POS tag). Also called *observation likelihoods*.


“**Hidden**” means the exact state sequence that generated the observations is not known.

HMM Probabilities



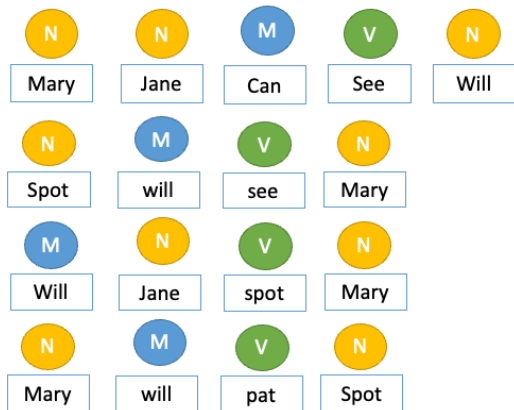
- HMM is a **generative model** that defines a joint probability distribution over the *observations* and *states*.
- Emission probabilities (i.e., Observation likelihoods) generate the *observations* given the *states*.

Hidden Markov Model for POS Tagging



- **Hidden Markov Model (HMM)** represented as *finite state machine*.
- Note that in this representation, the number of nodes (states) = the size of the set of POS tags.

Estimation of Emission Probabilities



Example sentences:

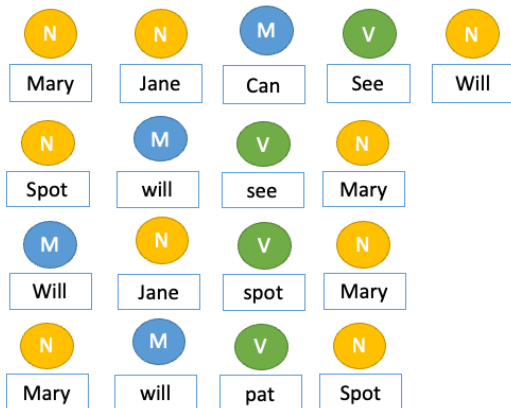
- *Mary Jane can see Will*
- *Spot will see Mary*
- *Will Jane spot Mary?*
- *Mary will pat Spot*

How to get Emission Probability?

Estimating of Emission Probabilities

Step 1: *Frequency count*

Word	N	M	V
Mary	4	0	0
Jane	2	0	0
Will	1	3	0
Spot	2	0	1
Can	0	1	0
See	0	0	2
pat	0	0	1
	9	4	4



Estimating of Emission Probabilities

Step 2: Normalization (column wise) → emission probability

Word	N	M	V
Mary	4	0	0
Jane	2	0	0
Will	1	3	0
Spot	2	0	1
Can	0	1	0
See	0	0	2
pat	0	0	1

Table: Raw Counts

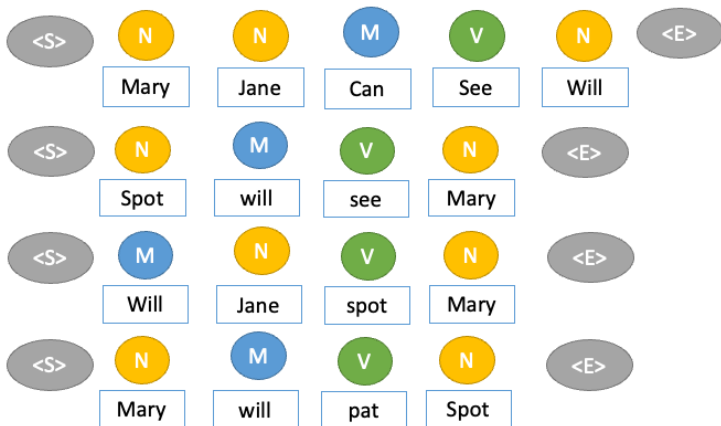
Word	N	M	V
Mary	4/9	0.00	0.00
Jane	2/9	0.00	0.00
Will	1/9	3/4	0.00
Spot	2/9	0.00	1/4
Can	0.00	1/4	0.00
See	0.00	0.00	2/4
pat	0.00	0.00	1/4

Table: Normalized

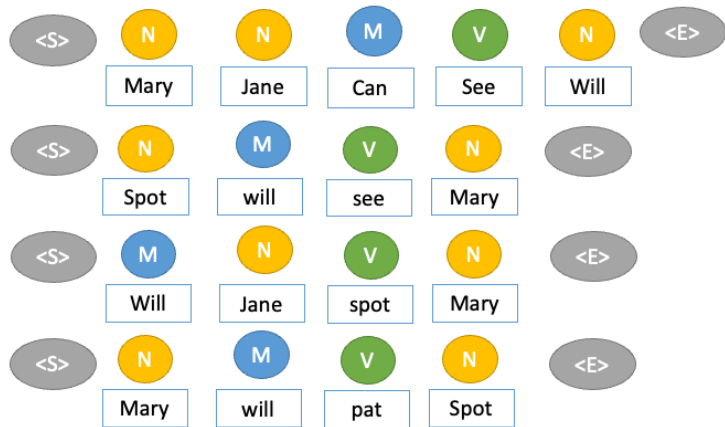
Estimating Transition Probabilities

Step 1: *introduce beginning and end tags*

- Probability of next POS tag given previous POS tag is called **transition probability**.
- To indicate *start* and *end* of tag, we introduce two symbols $\langle S \rangle$, $\langle E \rangle$ respectively.



Estimating Transition Probabilities



Step 2: Count co-occurrence of tags

	N	M	V	<E>
<S>	3	1	0	0
N	1	3	1	4
M	1	0	3	0
V	4	0	0	0

Estimating Transition Probabilities

Step 3: *Normalize (row wise)* → **Transition Probability**

	N	M	V	<E>
<S>	3	1	0	0
N	1	3	1	4
M	1	0	3	0
V	4	0	0	0


Table: Raw Counts

	N	M	V	<E>
<S>	3/4	1/4	0	0
N	1/9	3/9	1/9	4/9
M	1/4	0	3/4	0
V	4/4	0	0	0

Table: Normalized

Forward Algorithm

Determine the Likelihood of an Observation Sequence



Likelihood for this tag sequence:

$$3/4 \cdot 1/9 \cdot 3/9 \cdot 1/4 \cdot 3/4 \cdot 1/4 \cdot 1 \cdot 4/9 \cdot 4/9 = 0.0002572016$$

Viterbi Algorithm

HMM Decoding: Finds the *most likely sequence of states* (POS tags) that produced the observed sequence (words).


- Use transition and emission probability to predict next sequence.
- We don't want to search by enumerating all possible sequences of states. (Time complexity $O(N^T)$), where N is the number of states, T is the length of the sequence.
- **Dynamic Programming!**
- **Viterbi algorithm** has $O(N^2T)$ time complexity.

HMM Decoding: Viterbi Algorithm

Intuitions:


- The best path to a state at time t is the path that maximizes the probability to all the states at time $t - 1$ times the transition probability to the state at time t .
- So, it is possible that the best path to a state at time $t - 1$ is not the best path to the state at time t .
- We need a mechanism to find the best path to the final time step: **backtracing**.

HMM Tagging — Optimization



At each state, keep the **backpointer** to the previous state with the *highest probability*

HMM Tagging — Optimization



After the **forward computation**, a state will have only one incoming edge. Consequently, the **backtracing** will return the **best path of states** given the observation sequence.

function VITERBI(*observations* of len T , *state-graph* of len N) **returns** *best-path*

create a path probability matrix $viterbi[N+2, T]$

for each state s **from** 1 **to** N **do** ; initialization step

$viterbi[s, 1] \leftarrow a_{0,s} * b_s(o_1)$

$backpointer[s, 1] \leftarrow 0$

for each time step t **from** 2 **to** T **do** ; recursion step

for each state s **from** 1 **to** N **do**

$viterbi[s, t] \leftarrow \max_{s'=1}^N viterbi[s', t-1] * a_{s',s} * b_s(o_t)$

$backpointer[s, t] \leftarrow \operatorname{argmax}_{s'=1}^N viterbi[s', t-1] * a_{s',s}$

$viterbi[q_F, T] \leftarrow \max_{s=1}^N viterbi[s, T] * a_{s,q_F}$; termination step

$backpointer[q_F, T] \leftarrow \operatorname{argmax}_{s=1}^N viterbi[s, T] * a_{s,q_F}$; termination step

return the backtrace path by following backpointers to states back in time from $backpointer[q_F, T]$

Viterbi Algorithm

- 1 Problem Formalization
- 2 Part-of-Speech (POS) Tagging
- 3 Hidden Markov Model (HMM)
- 4 Named Entity Recognition (NER)**

3

³Slides adapted from CS447 UIUC <https://courses.grainger.illinois.edu/cs447/fa2020/Slides/Lecture24.pdf>

Name Entity Recognition (NER)

The task: find and classify names in text, for example:

The European Commission [ORG] said on Thursday it disagreed with German [MISC] advice. Only France [LOC] and Britain [LOC] backed Fischler [PER] 's proposal.

“What we have to be extremely careful of is how other countries are going to take Germany's lead”, Welsh National Farmers ' Union [ORG] (NFU [ORG]) chairman John Lloyd Jones [PER] said on BBC [ORG] radio.

Named Entity Types

Type	Tag	Sample Categories	Example sentences
<i>People</i>	PER	people, characters	Turing is a giant of computer science.
<i>Organization</i>	ORG	companies, sports teams	The IPCC warned about the cyclone.
<i>Location</i>	LOC	regions, mountains, seas	The Mt. Sanitas loop is in Sunshine Canyon .
<i>Geo-Political Entity</i>	GPE	countries, states, provinces	Palo Alto is raising the fees for parking.
<i>Facility</i>	FAC	bridges, buildings, airports	Consider the Golden Gate Bridge .
<i>Vehicles</i>	VEH	planes, trains, automobiles	It was a classic Ford Falcon .

These types were developed for the news domain as part of **NIST's Automatic Content Extraction (ACE) program**.

Other domains (e.g. *biomedical text*) require different types (proteins, genes, diseases, etc.)

NER on Word Sequences

NER on word sequences

We predict entities by *classifying words in context* and then extracting entities as **word sequences**.

Foreign	ORG	}	B-ORG
Ministry	ORG		I-ORG
spokesman	O		O
Shen	PER	}	B-PER
Guofang	PER		I-PER
told	O		O
Reuters	ORG	}	B-ORG
that	O		O
:	:		👉 BIO encoding

Why NER is Hard?

“**First National Bank Donates 2 Vans To Future School of Fort Smith**”

- **Hard to work out boundaries of entity**
 - ▶ Is the first entity First National Bank or National Bank?
- **Hard to know if something is an entity**
 - ▶ Is there a school Future School or is it a future school?
- **Hard to know class of unknown/novel entity:**
 - ▶ To find out more about Zig Ziglar and read features by other Creators Syndicate writers
 - ▶ What class is Zig Ziglar? (*A Person*)
- **Entity class is ambiguous and depends on context**
 - ▶ Where Larry Ellison and Charles Schwab can live discreetly amongst wooded estates.
 - ▶ Charles Schwab is *PER* not *ORG* here!

Sequence Labeling Algorithms for NER

Statistical models:

- Maximum Entropy Markov Models (MEMMs)
- Conditional Random Fields (CRFs)

Neural models:

- *Recurrent networks* (or *transformers*) that predict a label at each time step, possibly with a **CRF output layer**.

Maximum Entropy Markov Models

MEMMs use a *logistic regression* (Maximum Entropy) classifier for each $P(t^{(i)} \mid w^{(i)}, t^{(i-1)})$.

$$P(t^{(i)} = t_k \mid t^{(i-1)}, w^{(i)}) = \frac{\exp\left(\sum_j \lambda_{jk} f_j(t^{(i-1)}, w^{(i)})\right)}{\sum_l \exp\left(\sum_j \lambda_{jl} f_j(t^{(i-1)}, w^{(i)})\right)}$$

Here, $t^{(i)}$: label of the i -th word vs. t_i : i -th label in the inventory.

This requires the definition of a **feature function** $f(t^{(i-1)}, w^{(i)})$ that returns an n -dimensional feature vector for predicting label $t^{(i)} = t_j$ given inputs $t^{(i-1)}$ and $w^{(i)}$. Training returns weights λ_{jk} for each feature j used to predict label t_k .

Conditional Random Fields (CRFs)

Conditional Random Fields have the same mathematical definition as **MEMMs**, but:

- **CRFs** are trained *globally* to maximize the probability of the overall sequence,
- **MEMMs** are trained *locally* to maximize the probability of each individual label

This requires *dynamic programming*:

- **Training:** akin to the *Forward-Backward algorithm* used to train **HMMs** from unlabeled sequences
- **Decoding:** *Viterbi*

Feature-based NER (traditional approach)

identity of w_i , identity of neighboring words
embeddings for w_i , embeddings for neighboring words
part of speech of w_i , part of speech of neighboring words
base-phrase syntactic chunk label of w_i and neighboring words
presence of w_i in a **gazetteer**
 w_i contains a particular prefix (from all prefixes of length ≤ 4)
 w_i contains a particular suffix (from all suffixes of length ≤ 4)
 w_i is all upper case
word shape of w_i , word shape of neighboring words
short word shape of w_i , short word shape of neighboring words
presence of hyphen

- **Train a sequence labeling model** (MEMM or CRF), using features such as the ones listed above for English.

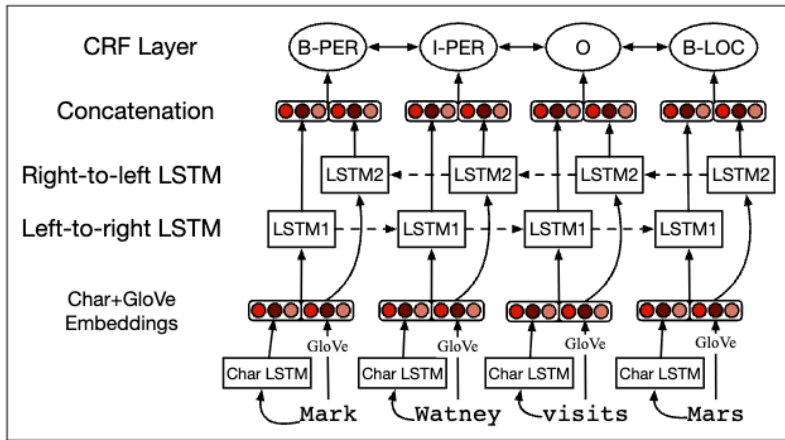
Feature-based NER (Cont.)

- *Word Shape*: replace all upper-case letters with one symbol (e.g. “X”), all lower-case letters with another symbol (“x”), all digits with another symbol (“d”), and leave punctuation marks as is (“L’Occitane → “X’Xxxxxxxxx”).
- *Short Word Shape*: remove adjacent letters that are identical in word shape (“L’Occitane → “X’Xxxxxxxxx” → “X’Xx”).

Neural NER

Sequence RNN (e.g. *biLSTM* or *Transformer*) with a **CRF output layer**.

Input: word embeddings, possibly concatenated with character embeddings and other features, e.g.:



Rule-based NER

The textbook gives an example of an *iterative approach* that makes multiple passes over the text:

- **Pass 1:** Use **high-precision rules** to label (a small number of) unambiguous mentions
- **Pass 2:** Propagate the labels of the previously detected named entities to any mentions that are substrings (or acronyms?) of these entities
- **Pass 3:** Use *application-specific name lists* to identify further likely names (as *features?*)
- **Pass 4:** Now use a *sequence labeling approach* for NER, keeping the already labeled entities as **high-precision anchors**.

The basic ideas behind this approach (*label propagation*, using **high-precision items** as anchors) can be useful for other tasks as well.

Summary

Summary